

Robotic Localization

Intro

Welcome to week four of the Estimation and Learning module in the Robotics specialization. This week, we'll look at localization algorithms for robotics. First of all, we'll explore how robots represent the pose, that is where they're located and how they can estimate this over time using a localization algorithm. Then, we'll see how, given a map, we can use the information in the map to help us build up this registration of the pose to the map. Finally, we'll explore a canonical algorithm called the particle filter which robots can use to estimate their localization information using sensors. In summary, this week you will learn how robots can keep track of where they are in space and time using these particle filter algorithms to represent their uncertainty.

4.1 Odometry Modeling

This week we will learn about self-localization techniques including the particle filter. In this first lecture, we will consider models for odometry as a first order approximation to the robot's location. As in your car, where the odometer records how many miles you have traveled, odometry provides a measurement of how far the robot has moved. Odometry is just one method of finding the robot's location in the world. If we look at a typical application of localization, car navigation, we see several ways to find location. Information sources include GPS, global positioning system, cellular networks, and Wi-Fi access points. Each of these sources, however, have certain levels of noise that lead to various levels of accuracy. Driverless cars, for instance, will need better than 3.5 meters of accuracy that the GPS provides. That error is the difference between occupying the sidewalk and the road. The previous sources represent global knowledge of position, exact coordinates. Odometry and other sources of information can augment the global localization sources with local knowledge. How have they changed coordinates? These sources of information are more precise, giving centimeter accuracy. However, integrating sources, like encoders and gyroscopes, over time can lead to drift. This is due to the accumulation of errors in time. Errors from slippage of the wheels deceive the encoder for instance. Other local sensors like laser scanners and color and depth cameras can help

to correct these errors. We will see how this incorporation happens later in the week. Odometry updates start with modeling the robot. Different robots, such as humanoids or aerial vehicles, will require different models. In our case, we will model a skid steer four-wheeled robot. The odometry measurements come from ticks from the encoder that measure how much the wheels have rotated in a given timeframe. These ticks can be mapped into translation and rotation of the body of the robot. First, let's explore the rotation odometry calculation. With a skid steer robot, the left and right sets of wheels are controlled independently. When turning, these two sides form the inner and outer radii of circles that share the same center. Coupled with the knowledge of how wide the robot vehicle is, we can determine a change in angle based on these encoder ticks. First, we want to translate motor ticks into meters traveled by the inner and outer wheels along their respective arcs. This conversion requires knowledge of the wheel sizes. Here, these measurements in meters are denoted e_o and e_i . The inner and outer arcs are known, but they also share the same angle of rotation. With knowledge of the width of the robot, we can use the difference in arcments to calculate the shared angle data. Next, we will consider the translation of the robot. Conveniently, the translation requires knowledge of the rotation that we have already calculated. In measuring translation, we can form a triangle with the known angle of rotation. We then can average the change in position for both the inner and outer wheel sets to find the change in the x direction. The change in the y direction requires a similar methodology. For small movements, this is a good approximation for the translation. Unfortunately, the encoder measurements can be noisy due to wheel slippage. Angular estimates then propagate errors into the translation estimates. One solution to this problem is to utilize the gyroscope to find a more precise measurement of angular change. For a small number of time steps, the gyroscope can be very accurate. Thus, angular odometry is measured solely by the rate of change observed by the gyro, integrated over time. This measurement aids in translation calculations as well. This simple odometry approach to localization requires a frame of reference for where the robot began its trip. Local measurements from the encoders and

gyroscopes still provide noisy estimates. So we want to include more measurements to correct errors. The next sections will discuss using maps to aid in localization correction, as well as ways to probabilistically define our localization state.

4.2 Map Registration

In this lecture, we will consider correlation based matching strategies for location a robot on a map given laser range data. This map registration process provides a very precise complement to odometry based localization. First we should introduced the LIDAR depth sensor. LIDAR stands for light detection and ranging and it provides distance measurements. Often engineered in a laser scanner to provide two dimensional data. The laser scanner we will model in this lecture takes depth measurements in polar coordinates, where a continuous distance reading r is made at discrete angles θ . Here, θ encompasses 270 degrees, not a full circle. The laser scanner can only see 10 to 30 meters away. In this range restriction, means that distance measurements showing here is black dots, can only be found within the area in green. Thus, due to the rays generating from a single point and the limited range, the robot can only see the dotted lines and not the lines in brown. Just as in the previous lecture, a two dimensional occupancy grid map will be used in localization, where a light colored cell represents high probability of an obstacle and a dark colored cell present a low probability. The cells here are meant to replicate the laser skin shown in the previous slide when the robot is approaching a corner in a hallway. Because the robot lives in a finite grid world the grid must sometimes be expanded as the robot can escape the boundaries. In this case, the map representation should increase in size as the robot turns and travels on the corridor shown in the top left of this map or else information will be lost. In addition to mapping the laser data discussed in week 3, we can access map data and try to find the robot pose in the map given the laser data. The complimentary stages of mapping and localization when performed together are known as SLAM, simultaneous localization and mapping, which is a major

research topic in robotics. In the localization problem we have two sets of information. First, the occupancy grid map provides a ground truth knowledge of what the robot should expect to observe in the world. Second, the set of lighter scan measurements provides information on what the robot is observing at the current time. The lighter scan measurements must be discretized according to the map representation, as discussed in week three, in order to be compared to the information from the occupancy map. With these two pieces of information the goal is to find the best robot pose on the map that explains the measured observations. Searching over all possible poses of the robot can be difficult. But based on the odometry information discussed in the last lecture, we have some tricks to make the search easier. We can constrain the search to a limited number of poses based on odometry information. Because we track the robot over time, we have the last known position of the robot and odometry information on how far the robot most likely moved. Thus, the most likely pose for the robot is now given a new set of laser data, is probably close to where the odometry predicts the robot to be. This prediction means that we can refine our search to poses near the prediction and be more confident in the validity of our search results. We measure each pose p in the search based on a map registration metric. One metric is to consider the sum of the map values m , at coordinates x and y , where the laser returns r , hit. This correlation metric can be modified to suit the application at hand. In our case, the value of our map cell will be a log odds ratio, so laser returns that are seen at a map location with high probability of occupancy will strongly increase the registration in the metric score. Laser returns with map locations known as free cells will decrease the metric score. Additionally, the correlation can be scaled where returns from far distances affect the metric calculation less than nearby the laser returns. We register the robot on the map, at the pose that maximizes the registration metric. Thus, when the odometry is calculated, it uses this pose to predict a new position of the robot, in time. In addition to considering merely the laser returns, we can consider points for the laser returns penetrated. This calculation can further corroborate our map registration. It requires considerably more computation however. To capture pose uncertainty using

a simple Gaussian on position and angle may not provide a feasible approach. In the next lecture, we will present a pose filter that can capture bi-modal uncertainty and non-linear models and a computationally tractable way.

4.3 Particle Filter

In this lecture, we will talk about a probabilistic state estimation technique using a sampling-based distribution representation known as the Particle Filter. Instead of a fully defined function, the Particle Filter represents a distribution with a set of samples, referred to as particles. These particles represent the distribution. The statistics of the samples match the statistics of the distribution, such as the mean or standard deviation. However, they can be more complicated metrics as well. In this way, there are no parameters as were seen in the mean and covariances of the Gaussian models. Instead, a full population is tracked. In essence, the particle filter population represents a mixture of Gaussian distributions that we have seen in the first week. Here, the variance will go to 0. With 0 variance, the Gaussian distributions become Dirac Delta functions. Initially, a set of particles represent the underlying belief state. Each particle is a pair of the pose and the weight of that pose. This is similar to representing a probability function where the weight is the probability of that pose in the underlying distribution. Here, darker colors represent higher weights, and lighter colors represent lower weights. Just like the Kalman filter, a motion model will move the underlying distribution. Here, the particles move based on odometry measurements taken from the robot. A companion uncertainty model captures the noise underlying the motion model. For instance, this could be wheel slip or friction changes. In the particle filter, where we do not track the motion model in explicit parameters, we add sampled noise from the motion noise model. In this case, we use a Gaussian distribution to model noise with 0 mean and non-0 covariance. Noise is uniquely added to each particle. So separate samples are made for each particle. After the noise is added, the dispersion of the particles captures the uncertainty due to movement. Like the Kalman filter, we can use a separate set of observations to constrain our noise and update our belief distribution. Here we will leverage

the LIDAR correlation from previous lectures on map registration. We will update the weights of the particles to reflect the correlation score from the map registration by utilizing the current weights as a prior belief. The new set of particles captures the distribution after odometry and sensor measurements. However, this may not be the optimal set to represent the distribution. Here, you can see that only a few particles have significant weights. Most of the particles are lightly colored and do not give much information about the distribution. To make the set of particles more accurately represent the belief state distribution, we check the number of effective particles. The number of effective particles acts as a criterion for when to resample particles. This resampling process provides a probabilistically motivated way to prune out lower weighted particles. With the set of large and small weights, using the cumulative probability function can aid in sampling. With normalized weights, the sum of the weights is 1, and can be represented as a monotonically increasing cumulative function. We sample a number, uniformly, between 0 and 1 of the cumulative range and find which weight includes that number. The particles with the indices found in the resampling approach become the new set of particles to be fed into the next odometry update. Particles may be duplicated, but the odometry noise will differentiate these particles. This approach provides a good way to approach a multi-nodal belief state distribution and non-linear effects of your motion model.